

Body text so the floats have somewhere to land.

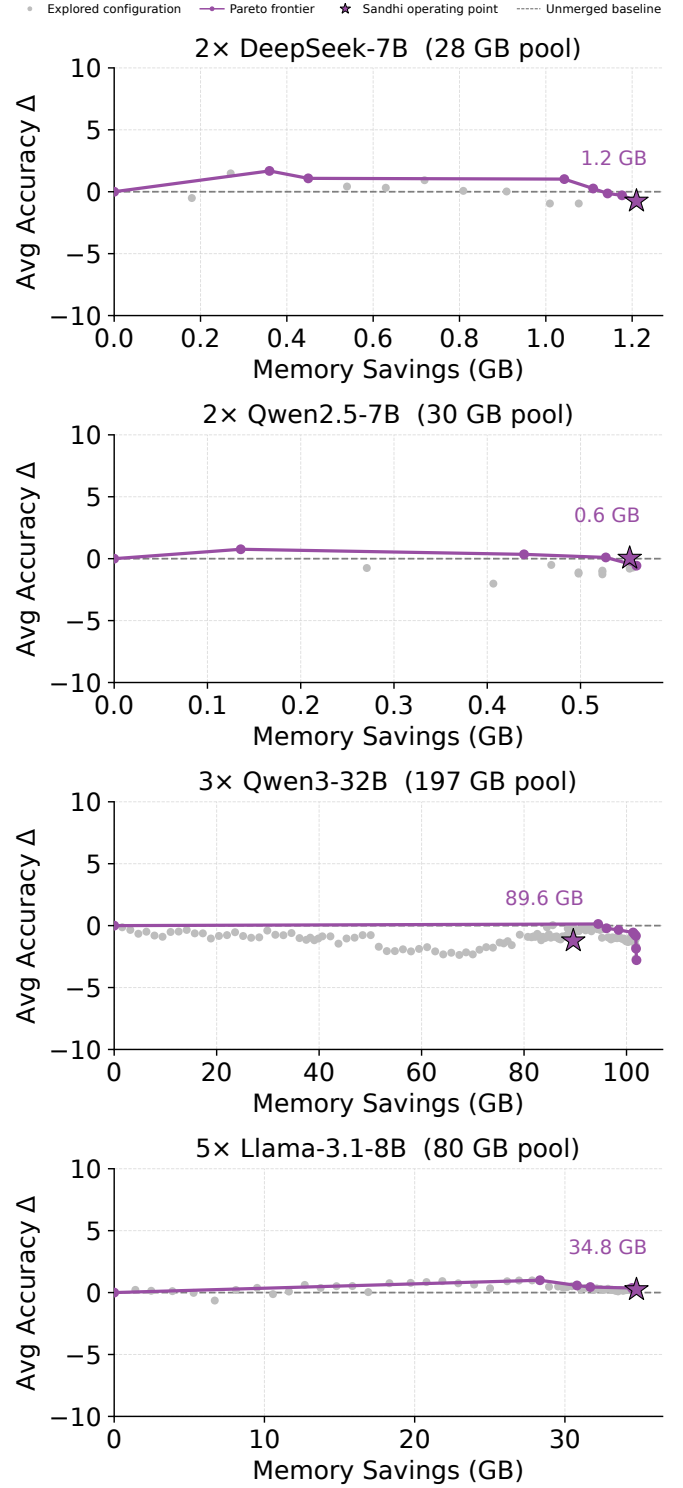


Figure 1: Accuracy–memory trade-off per model pool (average accuracy Δ across the pool’s models; continued in Figure 2). Grey points are explored merge configurations, the purple line is their Pareto frontier anchored at the unmerged pool, and the star is Sandhi’s selected operating point, chosen so *every* model stays within the user-set accuracy bound of its own unmerged baseline.



Figure 2: Accuracy-memory trade-off per model pool, continued from Figure 1: the 7-, 9-, and 12-model pools. The 9- and 12-model pools are additive compositions of the atom pools' sweeps (savings sum; accuracy is the model-count-weighted average), so every grey point is a deployable combination of per-pool configurations.